

```
In [1]: import pandas as pd
import numpy as np
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv("C:/Users/sahil/Documents/ecommerce.csv")
print(df)
```

	Order ID	Order Date	Ship Date	Aging	Ship Mode \
0	AU-2015-1	11/9/15	11/17/15	8.0	First Class
1	AU-2015-2	6/30/15	7/2/15	2.0	First Class
2	AU-2015-3	12/5/15	12/13/15	8.0	First Class
3	AU-2015-4	5/9/15	5/16/15	7.0	First Class
4	AU-2015-5	7/9/15	7/18/15	9.0	First Class
...
51285	FA-2015-30771	1/21/15	1/27/15	6.0	Standard Class
51286	FA-2015-30772	6/22/15	6/24/15	2.0	Standard Class
51287	FA-2015-30773	1/1/15	1/7/15	6.0	Standard Class
51288	FA-2015-30774	12/7/15	12/14/15	7.0	Standard Class
51289	FA-2015-30775	12/1/15	12/6/15	5.0	Standard Class

	Product Category	Product	Sales	Quantity	Discount	...	\
0	Auto & Accessories	Car Media Players	\$140.00	2	0.05	...	
1	Auto & Accessories	Car Speakers	\$211.00	3	0.03	...	
2	Auto & Accessories	Car Body Covers	\$117.00	5	0.01	...	
3	Auto & Accessories	Car & Bike Care	\$118.00	2	0.05	...	
4	Auto & Accessories	Tyre	\$250.00	1	0.04	...	
...	
51285	Fashion	Sports Wear	\$85.00	5	0.04	...	
51286	Fashion	Sports Wear	\$85.00	1	0.03	...	
51287	Fashion	Sports Wear	\$85.00	1	0.05	...	
51288	Fashion	Sports Wear	\$85.00	3	0.04	...	
51289	Fashion	Sports Wear	\$85.00	3	0.03	...	

	Shipping Cost	Order Priority	Customer ID	Customer Name	Segment	\
0	\$4.60	Medium	LS-001	Lane Daniels	Consumer	
1	\$11.20	Medium	IZ-002	Alvarado Kriz	Home Office	
2	\$3.10	Critical	EN-003	Moon Weien	Consumer	
3	\$2.60	High	AN-004	Sanchez Bergman	Corporate	
4	\$16.00	Critical	ON-005	Rowe Jackson	Corporate	
...	
51285	\$1.70	Medium	IN-0040977	Welch Fein	Corporate	
51286	\$0.20	Medium	TT-0040978	Martinez Arnett	Corporate	
51287	\$0.10	Medium	ON-0040979	Mccoy Duston	Home Office	
51288	\$2.80	Medium	RN-0040980	Bentley Zypern	Consumer	
51289	\$2.80	Medium	RZ-0040981	Mcclure Schwarz	Home Office	

	City	State	Country	Region	Months
0	Brisbane	Queensland	Australia	Oceania	Nov
1	Berlin	Berlin	Germany	Central	Jun
2	Porirua	Wellington	New Zealand	Oceania	Dec
3	Kabul	Kabul	Afghanistan	Central Asia	May
4	Townsville	Queensland	Australia	Oceania	Jul
...
51285	Pasadena	Texas	United States	Central	Jan
51286	Harare	Harare	Zimbabwe	Africa	Jun
51287	Townsville	Queensland	Australia	Oceania	Jan
51288	Houston	Texas	United States	Central	Dec
51289	Valinhos	São Paulo	Brazil	South	Dec

[51290 rows x 21 columns]

```
C:\Users\sahil\AppData\Local\Temp\ipykernel_18252\3304675805.py:1: DtypeWarning: Columns (8,9) have mixed types. Specify dtype option on import or set low_memory=False.  
df = pd.read_csv("C:/Users/sahil/Documents/ecommerce.csv")
```

```
In [3]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 51290 entries, 0 to 51289  
Data columns (total 21 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   Order ID              51290 non-null  object  
1   Order Date            51289 non-null  object  
2   Ship Date             51290 non-null  object  
3   Aging                 51289 non-null  float64  
4   Ship Mode             51290 non-null  object  
5   Product Category     51290 non-null  object  
6   Product               51290 non-null  object  
7   Sales                 51290 non-null  object  
8   Quantity              51289 non-null  object  
9   Discount              51290 non-null  object  
10  Profit                51290 non-null  object  
11  Shipping Cost         51290 non-null  object  
12  Order Priority         51288 non-null  object  
13  Customer ID           51289 non-null  object  
14  Customer Name         51290 non-null  object  
15  Segment               51289 non-null  object  
16  City                  51290 non-null  object  
17  State                 51290 non-null  object  
18  Country               51290 non-null  object  
19  Region                51289 non-null  object  
20  Months                51290 non-null  object  
dtypes: float64(1), object(20)  
memory usage: 8.2+ MB  
None
```

```
In [4]: print(df.isnull().sum())
```

Order ID	0
Order Date	1
Ship Date	0
Aging	1
Ship Mode	0
Product Category	0
Product	0
Sales	0
Quantity	1
Discount	0
Profit	0
Shipping Cost	0
Order Priority	2
Customer ID	1
Customer Name	0
Segment	1
City	0
State	0
Country	0
Region	1
Months	0

```
dtype: int64
```

```
In [8]: df[['Ship Date', 'Order Date']] = df[['Ship Date', 'Order Date']].apply(lambda col: p
```

[illegible]

```

-----
ValueError                                Traceback (most recent call last)
Cell In[9], line 2
      1 df[['Sales', 'Shipping Cost', 'Profit']] = (df[['Sales', 'Shipping Cost', 'P
rofit']].replace({r'\$': ''}, regex=True)
----> 2                                     .astype(float))

File ~\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\generi
c.py:6643, in NDFrame.astype(self, dtype, copy, errors)
    6637     results = [
    6638         ser.astype(dtype, copy=copy, errors=errors) for _, ser in self.items
    6639     ]
    6641 else:
    6642     # else, only a single dtype is given
-> 6643     new_data = self._mgr.astype(dtype=dtype, copy=copy, errors=errors)
    6644     res = self._constructor_from_mgr(new_data, axes=new_data.axes)
    6645     return res.__finalize__(self, method="astype")

File ~\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\interna
ls\managers.py:430, in BaseBlockManager.astype(self, dtype, copy, errors)
    427 elif using_copy_on_write():
    428     copy = False
--> 430 return self.apply(
    431     "astype",
    432     dtype=dtype,
    433     copy=copy,
    434     errors=errors,
    435     using_cow=using_copy_on_write(),
    436 )

File ~\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\interna
ls\managers.py:363, in BaseBlockManager.apply(self, f, align_keys, **kwargs)
    361     applied = b.apply(f, **kwargs)
    362     else:
--> 363     applied = getattr(b, f)(**kwargs)
    364     result_blocks = extend_blocks(applied, result_blocks)
    366 out = type(self).from_blocks(result_blocks, self.axes)

File ~\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\interna
ls\blocks.py:758, in Block.astype(self, dtype, copy, errors, using_cow, squeeze)
    755     raise ValueError("Can not squeeze with more than one column.")
    756     values = values[0, :] # type: ignore[call-overload]
--> 758 new_values = astype_array_safe(values, dtype, copy=copy, errors=errors)
    760 new_values = maybe_coerce_values(new_values)
    762 refs = None

File ~\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\dtypes
\astype.py:237, in astype_array_safe(values, dtype, copy, errors)
    234     dtype = dtype.numpy_dtype
    236 try:
--> 237     new_values = astype_array(values, dtype, copy=copy)
    238 except (ValueError, TypeError):
    239     # e.g. _astype_nansafe can fail on object-dtype of strings
    240     # trying to convert to float
    241     if errors == "ignore":

```

```
File ~\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\dtypes
\astype.py:182, in astype_array(values, dtype, copy)
    179     values = values.astype(dtype, copy=copy)
    181 else:
--> 182     values = _astype_nansafe(values, dtype, copy=copy)
    184 # in pandas we don't store numpy str dtypes, so convert to object
    185 if isinstance(dtype, np.dtype) and issubclass(values.dtype.type, str):

File ~\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\dtypes
\astype.py:133, in _astype_nansafe(arr, dtype, copy, skipna)
    129     raise ValueError(msg)
    131 if copy or arr.dtype == object or dtype == object:
    132     # Explicit copy, or required since NumPy can't view from / to object.
--> 133     return arr.astype(dtype, copy=True)
    135 return arr.astype(dtype, copy=copy)

ValueError: could not convert string to float: '0.xf'
```

```
In [10]: df['Sales'].value_counts()
```

```
Out[10]: Sales
$228.00    3823
$85.00     2827
$159.00    2796
$224.00    2795
$213.00    2795
$122.00    2795
$109.00    2795
$62.00     2795
$248.00    2794
$196.00    2794
$218.00    2794
$211.00    1853
$250.00    1114
$133.00    1053
$70.00     1029
$119.00    1029
$124.00    1029
$67.00     1029
$78.00     1029
$34.00     1028
$216.00    1027
$231.00     829
$114.00     827
$140.00     826
$54.00      826
$72.00      826
$117.00     826
$118.00     826
$130.00     261
$192.00     224
$83.00      224
$65.00      224
$199.00     221
$33.00      221
$104.00     221
$220.00     221
$111.00     221
$222.00     221
$149.00     221
0.xf         1
Name: count, dtype: int64
```

```
In [11]: df['Sales'] = df['Sales'].replace('0.xf', np.nan)
```

```
In [12]: df = (df
            .replace({'Sales': {'0.xf': np.nan},
                      'Shipping Cost': {'test': np.nan},
                      'Region': {'So3th': 'South', '4orth': 'North'},
                      'Quantity': {'abc': np.nan},
                      'Discount': {'xxx': np.nan}
                     }))
```

```
In [13]: data = (data
              .replace({'Sales': {'0.xf': np.nan},
```

```

        'Shipping Cost': {'test': np.nan},
        'Region': {'So3th': 'South', '4orth': 'North'},
        'Quantity': {'abc': np.nan}}})
    .fillna({'Order Date': '2015-04-17'})

```

NameError Traceback (most recent call last)

Cell In[13], line 1

```

----> 1 data = (data
      2         .replace({'Sales': {'0.xf': np.nan},
      3                     'Shipping Cost': {'test': np.nan},
      4                     'Region': {'So3th': 'South', '4orth': 'North'},
      5                     'Quantity': {'abc': np.nan}}})
      6         .fillna({'Order Date': '2015-04-17'}))

```

NameError: name 'data' is not defined

```

In [14]: df[['Sales', 'Shipping Cost', 'Profit']] = (df[['Sales', 'Shipping Cost', 'Profit']]
        .astype(float))
df.head(2)

```

Out[14]:

	Order ID	Order Date	Ship Date	Aging	Ship Mode	Product Category	Product	Sales	Quantity	Discount	...
0	AU-2015-1	2015-11-09	2015-11-17	8.0	First Class	Auto & Accessories	Car Media Players	140.0	2	0.05	...
1	AU-2015-2	2015-06-30	2015-07-02	2.0	First Class	Auto & Accessories	Car Speakers	211.0	3	0.03	...

2 rows × 21 columns

```

In [15]: df[df.columns[0:4]].head(5)

```

Out[15]:

	Order ID	Order Date	Ship Date	Aging
0	AU-2015-1	2015-11-09	2015-11-17	8.0
1	AU-2015-2	2015-06-30	2015-07-02	2.0
2	AU-2015-3	2015-12-05	2015-12-13	8.0
3	AU-2015-4	2015-05-09	2015-05-16	7.0
4	AU-2015-5	2015-07-09	2015-07-18	9.0

```

In [16]: df[['Quantity', 'Discount']] = df[['Quantity', 'Discount']].astype(float)
df.info()

```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51290 entries, 0 to 51289
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Order ID              51290 non-null  object
1   Order Date            51289 non-null  datetime64[ns]
2   Ship Date             51290 non-null  datetime64[ns]
3   Aging                 51289 non-null  float64
4   Ship Mode             51290 non-null  object
5   Product Category     51290 non-null  object
6   Product               51290 non-null  object
7   Sales                 51289 non-null  float64
8   Quantity              51288 non-null  float64
9   Discount              51289 non-null  float64
10  Profit                51290 non-null  float64
11  Shipping Cost         51289 non-null  float64
12  Order Priority         51288 non-null  object
13  Customer ID           51289 non-null  object
14  Customer Name         51290 non-null  object
15  Segment               51289 non-null  object
16  City                  51290 non-null  object
17  State                 51290 non-null  object
18  Country               51290 non-null  object
19  Region                51289 non-null  object
20  Months                51290 non-null  object
dtypes: datetime64[ns](2), float64(6), object(13)
memory usage: 8.2+ MB

```

```
In [17]: df.isnull().sum()
```

```

Out[17]: Order ID          0
Order Date          1
Ship Date           0
Aging               1
Ship Mode           0
Product Category    0
Product             0
Sales               1
Quantity            2
Discount            1
Profit              0
Shipping Cost       1
Order Priority       2
Customer ID         1
Customer Name       0
Segment             1
City                0
State               0
Country             0
Region              1
Months              0
dtype: int64

```

```
In [18]: import pandas as pd
import numpy as np
```

```

from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

# Using .loc for label-based indexing
column_name = 'Sales'
row_label = 793
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

# Encode the 'Product' column
le = LabelEncoder()
df['Product'] = le.fit_transform(df['Product'])

# Create a validation set by setting known values to NaN
validation_indices = [1, 3] # Indices of known values to set as NaN for validation
df_validation = df.copy()
df_validation.loc[validation_indices, 'Sales'] = np.nan

# Separate features and target for the imputer
features = df[['Product', 'Quantity', 'Discount']]
target = df[['Sales']]

# Combine features and target for the imputer
data_for_imputer = pd.concat([features, target], axis=1)

# Initialize KNNImputer
imputer = KNNImputer(n_neighbors=5)

# Fit and transform the data
imputed_data = imputer.fit_transform(data_for_imputer)

# Update the dataframe with imputed values
df['Sales'] = imputed_data[:, -1]

# Create a new DataFrame with imputed values
df_imputed = df_validation.copy()
df_imputed['Sales'] = imputed_data[:, -1]

# Reverse the encoding for the 'Product' column
df['Product'] = le.inverse_transform(df['Product'])
df_imputed['Product'] = le.inverse_transform(df_imputed['Product'].astype(int))

# Convert the 'Product' column back to object type for both DataFrames
df['Product'] = df['Product'].astype(object)
df_imputed['Product'] = df_imputed['Product'].astype(object)

# Verify that 'Product' column is now correctly decoded
print("Data after converting 'Product' column to object:\n", df['Product'].head(5))
# print("Imputed DataFrame after converting 'Product' column to object:\n", df_impu

# Calculate the mean squared error and R-squared for the imputed values
actual_values = df.loc[validation_indices, 'Sales']

```

```

imputed_values = df_imputed.loc[validation_indices, 'Sales']
mse = mean_squared_error(actual_values, imputed_values)
r2 = r2_score(actual_values, imputed_values)

print("Accuracy of Model")
print(f'Mean Squared Error of the imputation: {mse}')
print(f'R-squared of the imputation: {r2}')

# Using .loc for label-based indexing
column_name = 'Sales'
row_label = 793
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

```

Value at column 'Sales' and row 793: nan
Data after converting 'Product' column to object:

```

0    Car Media Players
1      Car Speakers
2    Car Body Covers
3    Car & Bike Care
4              Tyre
Name: Product, dtype: object
Accuracy of Model
Mean Squared Error of the imputation: 0.0
R-squared of the imputation: 1.0

```

Value at column 'Sales' and row 793: 211.0

In [19]: `df.isnull().sum()`

```

Out[19]: Order ID          0
Order Date        1
Ship Date         0
Aging             1
Ship Mode         0
Product Category  0
Product           0
Sales             0
Quantity          2
Discount          1
Profit            0
Shipping Cost     1
Order Priority    2
Customer ID       1
Customer Name     0
Segment          1
City              0
State             0
Country           0
Region            1
Months            0
dtype: int64

```

```
In [20]: import pandas as pd
import numpy as np
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

# Using .loc for label-based indexing
column_name = 'Shipping Cost'
row_label = 535
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

# Encode the 'Product' column
le = LabelEncoder()
df['Product'] = le.fit_transform(df['Product'])

# Create a validation set by setting known values to NaN
validation_indices = [1, 3] # Indices of known values to set as NaN for validation
df_validation = df.copy()
df_validation.loc[validation_indices, 'Shipping Cost'] = np.nan

# Separate features and target for the imputer
features = df[['Product', 'Quantity', 'Discount']]
target = df[['Shipping Cost']]

# Combine features and target for the imputer
data_for_imputer = pd.concat([features, target], axis=1)

# Initialize KNNImputer
imputer = KNNImputer(n_neighbors=5)

# Fit and transform the data
imputed_data = imputer.fit_transform(data_for_imputer)

# Update the dataframe with imputed values
df['Shipping Cost'] = imputed_data[:, -1]

# Create a new DataFrame with imputed values
df_imputed = df_validation.copy()
df_imputed['Shipping Cost'] = imputed_data[:, -1]

# Reverse the encoding for the 'Product' column
df['Product'] = le.inverse_transform(df['Product'])
df_imputed['Product'] = le.inverse_transform(df_imputed['Product'].astype(int))

# Convert the 'Product' column back to object type for both DataFrames
df['Product'] = df['Product'].astype(object)
df_imputed['Product'] = df_imputed['Product'].astype(object)

# Verify that 'Product' column is now correctly decoded
print("Data after converting 'Product' column to object:\n", df['Product'].head(5))
# print("Imputed DataFrame after converting 'Product' column to object:\n", df_impu
```

```

# Calculate the mean squared error and R-squared for the imputed values
actual_values = df.loc[validation_indices, 'Shipping Cost']
imputed_values = df_imputed.loc[validation_indices, 'Shipping Cost']
mse = mean_squared_error(actual_values, imputed_values)
r2 = r2_score(actual_values, imputed_values)

print("Accuracy of Model")
print(f'Mean Squared Error of the imputation: {mse}')
print(f'R-squared of the imputation: {r2}')

# Using .loc for label-based indexing
column_name = 'Shipping Cost'
row_label = 535
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

```

Value at column 'Shipping Cost' and row 535: nan
Data after converting 'Product' column to object:

0	Car Media Players
1	Car Speakers
2	Car Body Covers
3	Car & Bike Care
4	Tyre

Name: Product, dtype: object
Accuracy of Model
Mean Squared Error of the imputation: 0.0
R-squared of the imputation: 1.0

Value at column 'Shipping Cost' and row 535: 15.0

```

In [21]: import pandas as pd
import numpy as np
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

# Using .loc for label-based indexing
column_name = 'Quantity'
row_label = 95
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

# Encode the 'Product' column
le = LabelEncoder()
df['Product'] = le.fit_transform(df['Product'])

# Create a validation set by setting known values to NaN
validation_indices = [1, 3] # Indices of known values to set as NaN for validation
df_validation = df.copy()
df_validation.loc[validation_indices, 'Quantity'] = np.nan

```

```

# Separate features and target for the imputer
features = df[['Product', 'Quantity', 'Discount']]
target = df[['Quantity']]

# Combine features and target for the imputer
data_for_imputer = pd.concat([features, target], axis=1)

# Initialize KNNImputer
imputer = KNNImputer(n_neighbors=5)

# Fit and transform the data
imputed_data = imputer.fit_transform(data_for_imputer)

# Update the dataframe with imputed values
df['Quantity'] = imputed_data[:, -1]

# Create a new DataFrame with imputed values
df_imputed = df_validation.copy()
df_imputed['Quantity'] = imputed_data[:, -1]

# Reverse the encoding for the 'Product' column
df['Product'] = le.inverse_transform(df['Product'])
df_imputed['Product'] = le.inverse_transform(df_imputed['Product'].astype(int))

# Convert the 'Product' column back to object type for both DataFrames
df['Product'] = df['Product'].astype(object)
df_imputed['Product'] = df_imputed['Product'].astype(object)

# Verify that 'Product' column is now correctly decoded
print("Data after converting 'Product' column to object:\n", df['Product'].head(5))
# print("Imputed DataFrame after converting 'Product' column to object:\n", df_impu

# Calculate the mean squared error and R-squared for the imputed values
actual_values = df.loc[validation_indices, 'Quantity']
imputed_values = df_imputed.loc[validation_indices, 'Quantity']
mse = mean_squared_error(actual_values, imputed_values)
r2 = r2_score(actual_values, imputed_values)

print("Accuracy of Model")
print(f'Mean Squared Error of the imputation: {mse}')
print(f'R-squared of the imputation: {r2}')

# Using .loc for label-based indexing
column_name = 'Quantity'
row_label = 95
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

```

Value at column 'Quantity' and row 95: nan
 Data after converting 'Product' column to object:
 0 Car Media Players
 1 Car Speakers
 2 Car Body Covers
 3 Car & Bike Care
 4 Tyre
 Name: Product, dtype: object
 Accuracy of Model
 Mean Squared Error of the imputation: 0.0
 R-squared of the imputation: 1.0

Value at column 'Quantity' and row 95: 3.6

```
In [22]: import pandas as pd
import numpy as np
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

# Using .loc for label-based indexing
column_name = 'Discount'
row_label = 211
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

# Encode the 'Product' column
le = LabelEncoder()
df['Product'] = le.fit_transform(df['Product'])

# Create a validation set by setting known values to NaN
validation_indices = [1, 3] # Indices of known values to set as NaN for validation
df_validation = df.copy()
df_validation.loc[validation_indices, 'Discount'] = np.nan

# Separate features and target for the imputer
features = df[['Product', 'Quantity', 'Shipping Cost']]
target = df[['Discount']]

# Combine features and target for the imputer
data_for_imputer = pd.concat([features, target], axis=1)

# Initialize KNNImputer
imputer = KNNImputer(n_neighbors=5)

# Fit and transform the data
imputed_data = imputer.fit_transform(data_for_imputer)

# Update the dataframe with imputed values
df['Discount'] = imputed_data[:, -1]

# Create a new DataFrame with imputed values
df_imputed = df_validation.copy()
df_imputed['Discount'] = imputed_data[:, -1]
```

```

# Reverse the encoding for the 'Product' column
df['Product'] = le.inverse_transform(df['Product'])
df_imputed['Product'] = le.inverse_transform(df_imputed['Product'].astype(int))

# Convert the 'Product' column back to object type for both DataFrames
df['Product'] = df['Product'].astype(object)
df_imputed['Product'] = df_imputed['Product'].astype(object)

# Verify that 'Product' column is now correctly decoded
print("Data after converting 'Product' column to object:\n", df['Product'].head(5))
# print("Imputed DataFrame after converting 'Product' column to object:\n", df_impu

# Calculate the mean squared error and R-squared for the imputed values
actual_values = df.loc[validation_indices, 'Discount']
imputed_values = df_imputed.loc[validation_indices, 'Discount']
mse = mean_squared_error(actual_values, imputed_values)
r2 = r2_score(actual_values, imputed_values)

print("Accuracy of Model")
print(f'Mean Squared Error of the imputation: {mse}')
print(f'R-squared of the imputation: {r2}')

# Using .loc for label-based indexing
column_name = 'Discount'
row_label = 211
value_loc = df.loc[row_label, column_name]
print(f"\nValue at column '{column_name}' and row {row_label}: {value_loc}")

```

Value at column 'Discount' and row 211: nan

Data after converting 'Product' column to object:

0 Car Media Players

1 Car Speakers

2 Car Body Covers

3 Car & Bike Care

4 Tyre

Name: Product, dtype: object

Accuracy of Model

Mean Squared Error of the imputation: 0.0

R-squared of the imputation: 1.0

Value at column 'Discount' and row 211: 0.03

In [23]: `df.isna().sum()`


```
Out[23]: Order ID      0
Order Date    1
Ship Date     0
Aging         1
Ship Mode     0
Product Category 0
Product       0
Sales         0
Quantity      0
Discount      0
Profit        0
Shipping Cost 0
Order Priority 2
Customer ID   1
Customer Name 0
Segment       1
City          0
State         0
Country       0
Region        1
Months        0
dtype: int64
```

```
In [24]: df[df['Region'].isna()]
```

```
Out[24]:
```

	Order ID	Order Date	Ship Date	Aging	Ship Mode	Product Category	Product	Sales	Quantity	Discount
117	AU-2015-118	2015-08-16	2015-08-17	1.0	First Class	Auto & Accessories	Car Media Players	140.0	1.0	0.04

1 rows × 21 columns

```
In [25]: df[df.isna().any(axis=1)]
```

Out[25]:

	Order ID	Order Date	Ship Date	Aging	Ship Mode	Product Category	Product	Sales	Quantity	Discount
27	AU-2015-28	2015-09-29	2015-10-05	NaN	First Class	Auto & Accessories	Car Media Players	140.0	1.0	0.03
100	AU-2015-101	NaT	2015-04-23	6.0	First Class	Auto & Accessories	Car Speakers	211.0	5.0	0.03
117	AU-2015-118	2015-08-16	2015-08-17	1.0	First Class	Auto & Accessories	Car Media Players	140.0	1.0	0.04
131	AU-2015-132	2015-08-12	2015-08-15	3.0	First Class	Auto & Accessories	Bike Tyres	72.0	5.0	0.05
370	AU-2015-371	2015-11-21	2015-11-25	4.0	First Class	Auto & Accessories	Car Speakers	211.0	4.0	0.02
625	AU-2015-626	2015-02-03	2015-02-05	2.0	First Class	Auto & Accessories	Tyre	250.0	4.0	0.03
791	AU-2015-792	2015-04-27	2015-05-01	4.0	First Class	Auto & Accessories	Car Pillow & Neck Rest	231.0	1.0	0.01

7 rows × 21 columns

```
In [26]: df.loc[(df['Country'] == 'Italy') & (df['Region'].isna()), 'Region'] = 'South'
df.loc[117]
```

```
Out[26]: Order ID          AU-2015-118
Order Date      2015-08-16 00:00:00
Ship Date       2015-08-17 00:00:00
Aging           1.0
Ship Mode       First Class
Product Category Auto & Accessories
Product         Car Media Players
Sales           140.0
Quantity        1.0
Discount        0.04
Profit          54.4
Shipping Cost    5.4
Order Priority   High
Customer ID     RI-00118
Customer Name    Ayala Molinari
Segment         Consumer
City            Turin
State           Piedmont
Country         Italy
Region          South
Months          Aug
Name: 117, dtype: object
```

```
In [27]: df[df['Aging'].isna()]
```

```
Out[27]:
```

	Order ID	Order Date	Ship Date	Aging	Ship Mode	Product Category	Product	Sales	Quantity	Discount	..
27	AU-2015-28	2015-09-29	2015-10-05	NaN	First Class	Auto & Accessories	Car Media Players	140.0	1.0	0.03	.

1 rows × 21 columns

```
In [28]: df.loc[(df['Aging'].isna()), 'Aging'] = (df['Ship Date'] - df['Order Date']).dt.day
print(df.loc[27])
```

Order ID	AU-2015-28
Order Date	2015-09-29 00:00:00
Ship Date	2015-10-05 00:00:00
Aging	6.0
Ship Mode	First Class
Product Category	Auto & Accessories
Product	Car Media Players
Sales	140.0
Quantity	1.0
Discount	0.03
Profit	55.8
Shipping Cost	5.6
Order Priority	High
Customer ID	NG-0028
Customer Name	Harris Armstrong
Segment	Corporate
City	Jinan
State	Shandong
Country	China
Region	North Asia
Months	Sep

Name: 27, dtype: object

```
In [29]: df.loc[27, 'Aging'] = np.nan
```

```
In [30]: df.loc[(df['Order Date'].isna()), 'Order Date'] = df['Ship Date'] - pd.to_timedelta(
df.loc[100]
```

```
Out[30]:
```

Order ID	AU-2015-101
Order Date	2015-04-17 00:00:00
Ship Date	2015-04-23 00:00:00
Aging	6.0
Ship Mode	First Class
Product Category	Auto & Accessories
Product	Car Speakers
Sales	211.0
Quantity	5.0
Discount	0.03
Profit	99.4
Shipping Cost	9.9
Order Priority	Critical
Customer ID	RN-00101
Customer Name	Cook Bern
Segment	Consumer
City	Amsterdam
State	North Holland
Country	Netherlands
Region	Central
Months	Apr

Name: 100, dtype: object

```
In [ ]:
```